Supplementary Information: Situation engineering: separating evidence applicability from availability in AI systems

Geunsik Lim
College of Information and Communication Engineering, Sungkyunkwan University, Suwon, South Korea
leemgs@g.skku.edu

Supplementary Note 1: Benchmark schema

Each JSONL item contains id, category, query time, question, claims, user region (when the question states one), required variables, gold action and gold answer. Claims contain text and auditable event metadata. This is version 0.2, in which every category is balanced so that no constant answer and no recency rule exceeds chance on the decisive variable (Methods); version 0.1 and all of its results are archived in the repository under `data/v0.1/` and `results/v0.1/`.

Supplementary Note 2: Complete aggregate results

Method	Accuracy	SCHR	ECE
SCQA	1.000	0.000	0.052
Lexical-RAG	0.543	0.457	0.157
Recency-RAG	0.500	0.500	0.230
Latest-mention	0.474	0.526	0.276

SCHR is the complement of exact accuracy in this single-answer controlled setting and is not interpreted as an independent outcome. ECE is computed from rule-assigned confidence with small seeded perturbations; it is retained only as a software regression smoke test and is not evidence of neural calibration.

Supplementary Note 3: Category results

category	Latest-mention	Lexical-RAG	Recency-RAG	SCQA
counterfactual	0.477	1.000	0.500	1.000
hidden_premise	0.250	0.250	0.250	1.000
modality	0.770	0.730	0.750	1.000
observer	0.517	0.483	0.500	1.000
scope	0.500	0.500	0.500	1.000
source_conflict	0.497	0.503	0.500	1.000
temporal	0.307	0.337	0.500	1.000

Supplementary Note 4: Ablations

ablation	accuracy	errors
none	1.000	0
time	0.929	300
source	0.929	300
modality	0.929	300
scope	0.929	300
observer	0.929	300
world	0.929	300
missing	0.929	300

Each ablation produces exactly 300 errors, all in the category whose variable was removed. In v0.1 the source and world ablations produced no errors, because in all 600 source-conflict items the official claim was the most recent, as was the hypothetical claim in all 600 counterfactual items; v0.2 makes the decisive claim the most recent in exactly half of these items.

Supplementary Note 5: Reproduction

From the repository root, run `python code/run_diagnostic.py` to regenerate the benchmark and the oracle-state baselines, category results and ablations (Figs. 2–4; requires the packages in `code/requirements-diagnostic.txt`), then `python code/make_figures.py` to redraw the figures and `python code/make_supplement_tables.py` to rewrite the tables in Supplementary Notes 2–4. Run `PYTHONPATH=code python code/run_experiments.py` to regenerate gold-, predicted- and corrupted-state outputs. Run `PYTHONPATH=code pytest -q code/tests` for implementation tests.

Supplementary Note 6: Situation-engineering design checklist

Supplementary Table 1 lists, for each component of a situation-engineered system, the questions its documentation should answer.

Supplementary Table 1 Minimum reporting checklist for situation-engineered AI systems.

Component	Questions that should be answered
Situation schema	Which actors, events, valid times, scopes, modalities, observers, worlds and uncertainty variables are represented?
Sensing	Which variables are supplied, retrieved or inferred? How is extraction accuracy measured independently of final answers?
Assembly and update	How are event coreference, supersession, conflict, validity intervals and belief revision handled?
Policy	Under which conditions does the system answer, retrieve, clarify, abstain or act?
Calibration	Does state uncertainty propagate into answer confidence and selective risk?
Observability	Can an auditor recover the active state, provenance, alternatives and policy rationale for each output?
Ablation	Does removing a state variable create the predicted category-specific failure rather than only a generic score reduction?
External validity	Are naturally occurring, multilingual, multimodal, adversarial and distribution-shift settings evaluated?
Governance	Are sensitive state variables minimized, access-controlled, correctable and audited for disparate effects?

Supplementary Note 7: Situation-engineering maturity model

Supplementary Table 2 defines five maturity levels, each with the evidence that would falsify a claim to have reached it.

Supplementary Table 2 Proposed SE0–SE4 maturity levels.

Level	Capability	Falsifiable evidence
SE0	Implicit situation inference only	No external situation representation or state-specific evaluation
SE1	Isolated situation tags	Measured extraction of individual variables such as time, location or modality
SE2	Explicit dynamic state	Provenance-linked state, update rules and temporal or scope consistency tests
SE3	Epistemic control	Missing-variable detection, competing hypotheses and answer/retrieve/clarify/abstain policy
SE4	Adaptive audited situation intelligence	Learned multimodal sensing, multi-agent belief tracking, shift monitoring and independent state audits

Supplementary Note 8: Recommended future benchmark modules

A natural-text extension should separately score event extraction, event coreference, temporal ordering, validity intervals, source and modality classification, scope resolution, observer belief, counterfactual frame selection, missing-variable detection, clarification utility, selective risk and claim-level situation consistency. Fixed-state policy tests and fixed-policy sensing tests are both needed to prevent final accuracy from obscuring the source of improvement.

Supplementary Note 9: Frontier-model evaluation status

The repository includes a frozen temporal SituatedQA smoke-test sample, a multi-provider evaluation harness and archived item-level responses from an audited Gemini-family pilot. On SituationCatch-Bench v0.2, three Gemini models completed all direct, structured and situation conditions on the 210-item stratified sample (30 items in each category), with no failed calls. The repository also archives a 12-item SituatedQA temporal check, a plain-text retrieval baseline and a search-then-answer tool-using agent baseline, as well as the v0.1 pilot and aborted v0.2 attempts, which enter no result. The integrity gate joins responses to the frozen item set and refuses to generate a manuscript table unless every requested model, condition and category is complete. A pre-registered replication (E6) repeats every condition with three open-weight models from Meta, DeepSeek and Qwen, and a pre-registered sensor study (E7) uses four LLMs as situation sensors; both are complete, audited and archived with the original and amended pre-registration (`paper/PREREGISTRATION_E6_E7.md`). Proprietary families beyond Gemini were dropped from E6 when the API credit ran out. A confirmatory revision still requires the full licensed natural-language split, proprietary families beyond Gemini, human annotation and stronger temporal-retrieval and verification-agent baselines under matched budgets.

Supplementary Note 10: Independent annotation protocol

The executable workflow creates separately randomized packets for at least three annotators. Annotators label action, answer, temporal state, modality, scope, source status, observer state and possible world without seeing method identity or another annotator's labels. Files are frozen before agreement calculation and adjudication. The scorer rejects missing cells and reports slot-level Fleiss' κ and unanimous agreement. Recruitment, ethics status, compensation, qualifications, exclusions and adjudication changes must accompany any reported human result. No such result is claimed here. A separate three-persona run validates the workflow with injected disagreements; its files are machine-marked synthetic and are not human-subject evidence.

Supplementary Note 11: Interfaces among engineering disciplines

Supplementary Table 3 specifies what each neighbouring discipline supplies to the situation layer and what it must receive back.

Supplementary Table 3 Recommended typed interfaces between situation engineering and neighbouring components.

Neighbouring discipline	Input to the situation layer	Output or requirement from the situation layer
Prompt engineering	Task, desired output, constraints and user intent	Required situation variables and clarification policy
Context / retrieval engineering	Candidate claims, passages, timestamps and source metadata	Missing-variable queries, applicability filters and freshness requirements
Memory engineering	Events, commitments, observations and provenance across sessions	Validity, supersession, observer/world tags and forgetting constraints
Tool engineering	Tool schema, preconditions, results, side effects and error states	Permitted calls, required confirmation and normalized state transition
Protocol engineering	Capability negotiation, authorization scopes and resource identifiers	Situation schema version, provenance envelope and uncertainty representation
Harness engineering	Workflow trace, checkpoints, permissions and recovery state	State invariants, action gate and answer/retrieve/clarify/abstain decision
Evaluation engineering	Test scenarios, judges, counterfactual perturbations and production traces	State-specific expected behaviour and decomposed failure labels
Observability engineering	Logs, lineage, timing and causal trace	Active-state snapshot, alternatives, unresolved variables and policy rationale
Serving engineering	Model availability, latency, token, memory and cost budgets	Minimum verification budget and graceful degradation policy
Security engineering	Trust zones, identities, privileges and threat signals	State-dependent authorization and response restrictions

Supplementary Note 12: Engineering anti-patterns

The taxonomy enables failure analysis that is more precise than attributing every problem to "the prompt" or "the model". Common anti-patterns include: (i) context dumping, in which more documents are added without resolving which state is active; (ii) memory accumulation, in which obsolete beliefs are retained without validity or supersession metadata; (iii) tool proliferation, in which capabilities are exposed without preconditions, typed errors or state effects; (iv) harness opacity, in which retries and intermediate actions are not connected to an auditable situation state; (v) evaluation collapse, in which only final answer accuracy is measured and sensing, state

update and policy failures are indistinguishable; and (vi) serving-induced semantic degradation, in which cost or latency optimizations silently reduce retrieval depth, truncate decisive evidence or skip verification. A situation-engineered system should make these trade-offs explicit and test whether degraded resource modes preserve calibrated abstention rather than confident error.

Supplementary Note 13: Minimum cross-layer reporting standard

A complete report should identify: the model and prompt; context and retrieval policy; memory write/read/forget rules; tool schemas and error semantics; protocol and authorization assumptions; harness orchestration and recovery; situation schema, update rules and action gate; evaluation data, judges and confidence intervals; observability and provenance records; and serving budgets, latency, cost and failure modes. Reporting only the model name and prompt is insufficient for agentic systems because performance is jointly determined by the model, environment and scaffolding.

Supplementary Note 14: Prospective external-validity protocol

This note is a design record, not a report of observations: it fixes the strata, targets and analysis before any external-validity data are collected. No outcome values are shown, because none have been measured.

Supplementary Table 4 Pre-specified external-validity design (no outcomes measured).

Stratum	Target	Split rule	Role
Natural, in-domain	400	no template crosses splits	primary
Natural, held-out domain	400	domain unseen in development	confirmatory
Author-disjoint semi-synthetic	400	disjoint generator authors	secondary

The planned minimum is 1,200 scenarios balanced across three domains, with no template paraphrases crossing splits. Power must be recomputed from pilot scenario-level variance; 1,200 is a collection target, not a guarantee of adequate power. The primary estimand is paired exact-accuracy difference (situation minus structured), clustered by scenario and reported separately by model family. Secondary endpoints are applicability-label F1, clarification utility, abstention precision, selective risk at fixed coverage, latency, provider-reported tokens and cost. Holm-adjusted category analyses are secondary. A result is confirmatory only if the pre-specified positive contrast replicates in at least two independent families and on the held-out natural stratum, while the situation condition is non-inferior to temporal RAG within a two-percentage-point margin. All null and adverse results remain reportable.

Generator and solver teams remain separated until predictions are frozen; annotators are blinded to condition; exclusions, prompt changes and failed calls are logged; and human recruitment, consent/ethics determination, compensation and adjudication are documented. The decision rules were dry-run only on simulated values stored in `results/prospective_expected_results.json`, which carries a `synthetic: true` flag so submission audits prevent their promotion to evidence; those values are deliberately not reported here.