

# Supplementary Information: Situation engineering is a missing layer for reliable artificial intelligence

Geunsik Lim

College of Information and Communication Engineering, Sungkyunkwan University,  
Suwon, South Korea

leemgs@g.skku.edu

## Supplementary Note 1: Benchmark schema

Each JSONL item contains id, category, query time, question, claims, required variables, gold action and gold answer. Claims contain text and auditable event metadata.

## Supplementary Note 2: Complete aggregate results

| Method         | Accuracy | SCHR  | ECE   |
|----------------|----------|-------|-------|
| SCQA           | 1.000    | 0.000 | 0.053 |
| Latest-mention | 0.857    | 0.143 | 0.183 |
| Recency-RAG    | 0.714    | 0.286 | 0.144 |
| Lexical-RAG    | 0.700    | 0.300 | 0.094 |

SCHR is the complement of exact accuracy in this single-answer controlled setting and is not interpreted as an independent outcome. ECE is computed from rule-assigned confidence with small seeded perturbations; it is retained only as a software regression smoke test and is not evidence of neural calibration.

### Supplementary Note 3: Category results

| category        | Latest-mention | Lexical-RAG | Recency-RAG | SCQA  |
|-----------------|----------------|-------------|-------------|-------|
| counterfactual  | 1.000          | 1.000       | 1.000       | 1.000 |
| hidden_premise  | 0.000          | 0.000       | 0.000       | 1.000 |
| modality        | 1.000          | 1.000       | 1.000       | 1.000 |
| observer        | 1.000          | 0.530       | 0.000       | 1.000 |
| scope           | 1.000          | 1.000       | 1.000       | 1.000 |
| source_conflict | 1.000          | 1.000       | 1.000       | 1.000 |
| temporal        | 1.000          | 0.372       | 1.000       | 1.000 |

### Supplementary Note 4: Ablations

| ablation | accuracy | errors |
|----------|----------|--------|
| none     | 1.000    | 0      |
| time     | 0.935    | 273    |
| source   | 1.000    | 0      |
| modality | 0.857    | 600    |
| scope    | 0.857    | 600    |
| observer | 0.857    | 600    |
| world    | 1.000    | 0      |
| missing  | 0.857    | 600    |

### Supplementary Note 5: Reproduction

From the repository root, run `PYTHONPATH=code python code/run_experiments.py`. This regenerates gold-, predicted- and corrupted-state outputs. Run `PYTHONPATH=code pytest -q code/tests` for implementation tests.

### Supplementary Note 6: Situation-engineering design checklist

### Supplementary Note 7: Situation-engineering maturity model

### Supplementary Note 8: Recommended future benchmark modules

A natural-text extension should separately score event extraction, event coreference, temporal ordering, validity intervals, source and modality classification, scope resolution, observer belief, counterfactual frame selection, missing-variable detection,

**Table 1** Minimum reporting checklist for situation-engineered AI systems.

| Component           | Questions that should be answered                                                                                         |
|---------------------|---------------------------------------------------------------------------------------------------------------------------|
| Situation schema    | Which actors, events, valid times, scopes, modalities, observers, worlds and uncertainty variables are represented?       |
| Sensing             | Which variables are supplied, retrieved or inferred? How is extraction accuracy measured independently of final answers?  |
| Assembly and update | How are event coreference, supersession, conflict, validity intervals and belief revision handled?                        |
| Policy              | Under which conditions does the system answer, retrieve, clarify, abstain or act?                                         |
| Calibration         | Does state uncertainty propagate into answer confidence and selective risk?                                               |
| Observability       | Can an auditor recover the active state, provenance, alternatives and policy rationale for each output?                   |
| Ablation            | Does removing a state variable create the predicted category-specific failure rather than only a generic score reduction? |
| External validity   | Are naturally occurring, multilingual, multimodal, adversarial and distribution-shift settings evaluated?                 |
| Governance          | Are sensitive state variables minimized, access-controlled, correctable and audited for disparate effects?                |

**Table 2** Proposed SE0–SE4 maturity levels.

| Level | Capability                              | Falsifiable evidence                                                                                   |
|-------|-----------------------------------------|--------------------------------------------------------------------------------------------------------|
| SE0   | Implicit situation inference only       | No external situation representation or state-specific evaluation                                      |
| SE1   | Isolated situation tags                 | Measured extraction of individual variables such as time, location or modality                         |
| SE2   | Explicit dynamic state                  | Provenance-linked state, update rules and temporal or scope consistency tests                          |
| SE3   | Epistemic control                       | Missing-variable detection, competing hypotheses and answer/retrieve/clarify/abstain policy            |
| SE4   | Adaptive audited situation intelligence | Learned multimodal sensing, multi-agent belief tracking, shift monitoring and independent state audits |

clarification utility, selective risk and claim-level situation consistency. Fixed-state policy tests and fixed-policy sensing tests are both needed to prevent final accuracy from obscuring the source of improvement.

## Supplementary Note 9: Frontier-model evaluation status

The repository includes a frozen temporal SituatedQA smoke-test sample, a multi-provider evaluation harness and archived item-level responses from an audited Gemini-family pilot. Three Gemini models completed all direct, structured and situation conditions on the 210-item stratified synthetic sample (30 items in each category), with no failed calls. The repository also archives a 12-item SituatedQA temporal check,

a top-3 lexical-retrieval baseline and a search-then-answer tool-using agent baseline. The integrity gate joins responses to the frozen item set and refuses to generate a manuscript table unless every requested model, condition and category is complete. Thus these are real, auditable results, but only a single-provider preliminary test. A confirmatory revision still requires the full licensed natural-language split, multiple independent model families and stronger temporal-retrieval and verification-agent baselines under matched budgets.

## **Supplementary Note 10: Independent annotation protocol**

The executable workflow creates separately randomized packets for at least three annotators. Annotators label action, answer, temporal state, modality, scope, source status, observer state and possible world without seeing method identity or another annotator’s labels. Files are frozen before agreement calculation and adjudication. The scorer rejects missing cells and reports slot-level Fleiss’  $\kappa$  and unanimous agreement. Recruitment, ethics status, compensation, qualifications, exclusions and adjudication changes must accompany any reported human result. No such result is claimed here. A separate three-persona run validates the workflow with injected disagreements; its files are machine-marked synthetic and are not human-subject evidence.

## **Supplementary Note 11: Interfaces among engineering disciplines**

## **Supplementary Note 12: Engineering anti-patterns**

The taxonomy enables failure analysis that is more precise than attributing every problem to “the prompt” or “the model”. Common anti-patterns include: (i) context dumping, in which more documents are added without resolving which state is active; (ii) memory accumulation, in which obsolete beliefs are retained without validity or supersession metadata; (iii) tool proliferation, in which capabilities are exposed without preconditions, typed errors or state effects; (iv) harness opacity, in which retries and intermediate actions are not connected to an auditable situation state; (v) evaluation collapse, in which only final answer accuracy is measured and sensing, state update and policy failures are indistinguishable; and (vi) serving-induced semantic degradation, in which cost or latency optimizations silently reduce retrieval depth, truncate decisive evidence or skip verification. A situation-engineered system should make these trade-offs explicit and test whether degraded resource modes preserve calibrated abstention rather than confident error.

## **Supplementary Note 13: Minimum cross-layer reporting standard**

A complete report should identify: the model and prompt; context and retrieval policy; memory write/read/forget rules; tool schemas and error semantics; protocol

**Table 3** Recommended typed interfaces between situation engineering and neighbouring components.

| Neighbouring discipline         | Input to the situation layer                                               | Output or requirement from the situation layer                                 |
|---------------------------------|----------------------------------------------------------------------------|--------------------------------------------------------------------------------|
| Prompt engineering              | Task, desired output, constraints and user intent                          | Required situation variables and clarification policy                          |
| Context / retrieval engineering | Candidate claims, passages, timestamps and source meta-data                | Missing-variable queries, applicability filters and freshness requirements     |
| Memory engineering              | Events, commitments, observations and provenance across sessions           | Validity, supersession, observer/world tags and forgetting constraints         |
| Tool engineering                | Tool schema, preconditions, results, side effects and error states         | Permitted calls, required confirmation and normalized state transition         |
| Protocol engineering            | Capability negotiation, authorization scopes and resource identifiers      | Situation schema version, provenance envelope and uncertainty representation   |
| Harness engineering             | Workflow trace, checkpoints, permissions and recovery state                | State invariants, action gate and answer/retrieve/clarify/abstain decision     |
| Evaluation engineering          | Test scenarios, judges, counterfactual perturbations and production traces | State-specific expected behaviour and decomposed failure labels                |
| Observability engineering       | Logs, lineage, timing and causal trace                                     | Active-state snapshot, alternatives, unresolved variables and policy rationale |
| Serving engineering             | Model availability, latency, token, memory and cost budgets                | Minimum verification budget and graceful degradation policy                    |
| Security engineering            | Trust zones, identities, privileges and threat signals                     | State-dependent authorization and response restrictions                        |

and authorization assumptions; harness orchestration and recovery; situation schema, update rules and action gate; evaluation data, judges and confidence intervals; observability and provenance records; and serving budgets, latency, cost and failure modes. Reporting only the model name and prompt is insufficient for agentic systems because performance is jointly determined by the model, environment and scaffolding.